

# QEMU 中 RISC-V SoC 的新增与实现

——基于NucLei SOC 在QEMU中的扩展

PLCT 王俊强

wangjunqiang@iscas.ac.cn

# 目录

| QEMU与RISC-V

| NucLei Soc简述


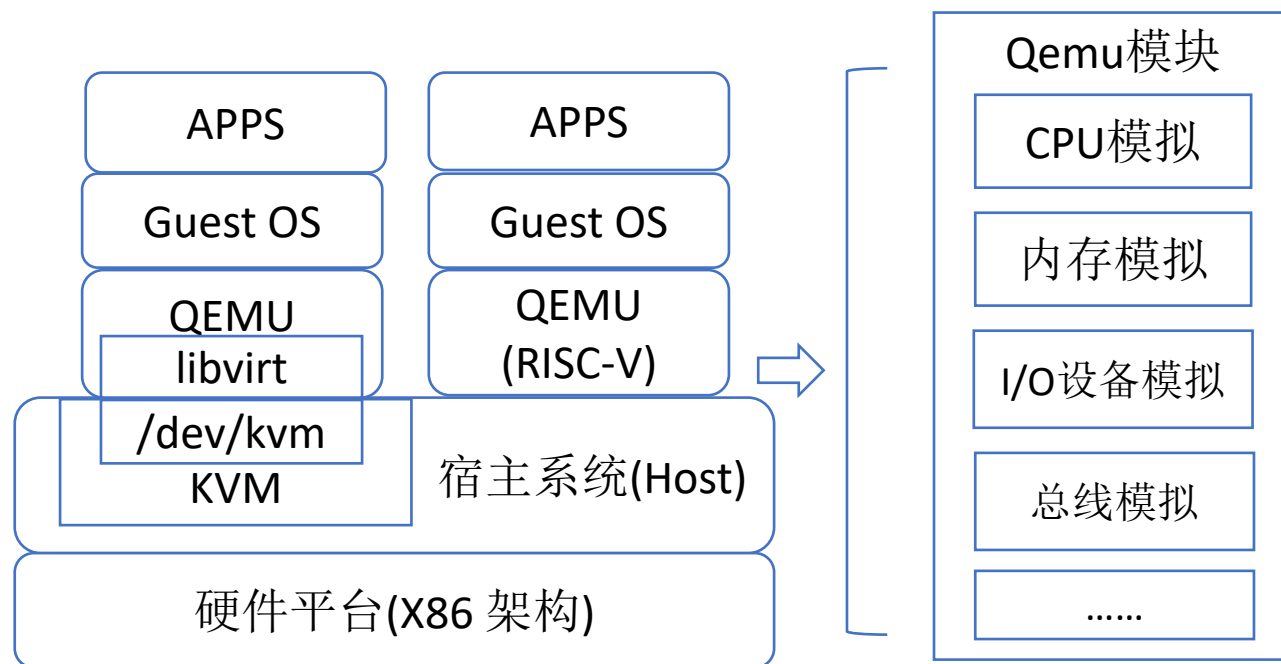
| QEMU中NucLei Soc的添加

| GD32V103F Board for RTOS

| HummingBird Board for Linux

# QEMU RISC-V

- QEMU是一款开源的**模拟器**及**虚拟机监管器**(Virtual Machine Monitor, VMM)
- RISC-V是一个基于精简指令集（RISC）原则的**开源指令集架构**（ISA）



Fedora Desktop for RISC-V

# QEMU RISC-V

## QEMU对RISC-V支持

Features	support
Hardware Virtualization	soon* <a href="#">kvm-riscv</a>
TCG Guest	Yes
MTTCG	Yes
TCG Host	No

Booting Linux	support
Debian	64
Fedora	64
OpenEmbedded	32/64
Buildroot	32/64

32bit Supported machines are: none empty machine opentitan RISC-V Board compatible with OpenTitan sifive_e RISC-V Board compatible with SiFive E SDK sifive_u RISC-V Board compatible with SiFive U SDK spike RISC-V Spike board (default) virt RISC-V VirtIO board	<b>cpu type:</b> any lowrisc-ibex rv32 sifive-e31 sifive-e34 sifive-u34
64bit Supported machines are: microchip-icicle-kit Microchip PolarFire SoC Icicle Kit none empty machine sifive_e RISC-V Board compatible with SiFive E SDK sifive_u RISC-V Board compatible with SiFive U SDK spike RISC-V Spike board (default) virt RISC-V VirtIO board	<b>cpu type:</b> any rv64 sifive-e51 sifive-u54

privileged 1.10.0  
privileged 1.11.0  
v extension 0.7.1

Extension	Version
I	2.1
E	2.1
M	2.0
A	2.1
F	2.2
D	2.2
V	0.7.1
C	2.0
H	0.3?
Counters	2.0
Zifencei	2.0
Zicsr	2.0

Extension支持(参考)

# NucLei Soc简述

N 级别			NX 级别			UX 级别			可选特性
32位架构 MCU, 边缘计算, AIoT, 安全			64位架构 存储, AR/VR, AI			64位架构 Linux, 数据中心, 网络设备, 基带			
900 系列 9-Stage Pipeline Dual-issue	N900	对标 ARM Cortex M7 ARM Cortex R4 ARM Cortex R5 ARM Cortex R7	NX900	NX900 多核	对标 ARM Cortex M7 ARM Cortex R5 ARM Cortex R7 ARM Cortex R8	UX900	UX900 多核	对标 ARM Cortex A9 ARM Cortex A53	
600 系列 6-Stage Pipeline Single-issue	N600	对标 ARM Cortex M7 ARM Cortex R4 ARM Cortex R5	NX600	NX600 多核	对标 ARM Cortex M7 ARM Cortex R4 ARM Cortex R5	UX600	UX600 多核	对标 ARM Cortex A5 ARM Cortex A7	
300 系列 3-Stage Pipeline Single-issue	N300	对标 ARM Cortex M33 ARM Cortex M4 ARM Cortex M4F							
200 系列 2-Stage Pipeline Single-issue	N200	对标 ARM Cortex M0 ARM Cortex M0+ ARM Cortex M3 ARM Cortex M23							
100 系列 2-Stage Pipeline Single-issue	N100	对标 8051 ARM Cortex M0 ARM Cortex M0+							

安全

可靠

扩展

DSP

浮点


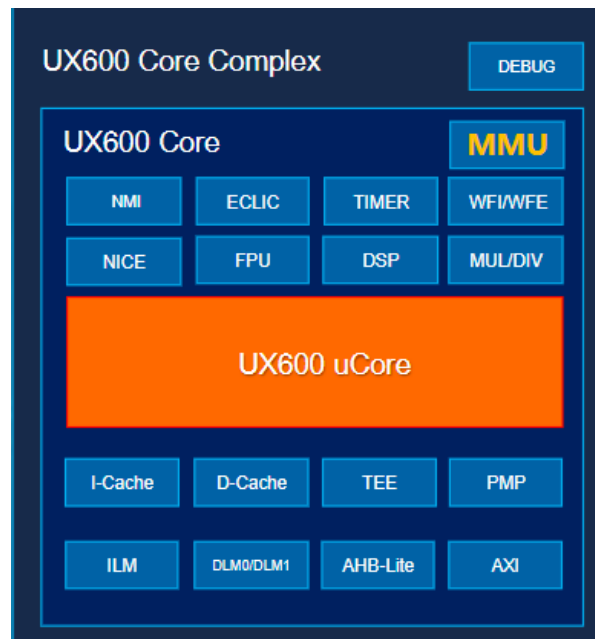
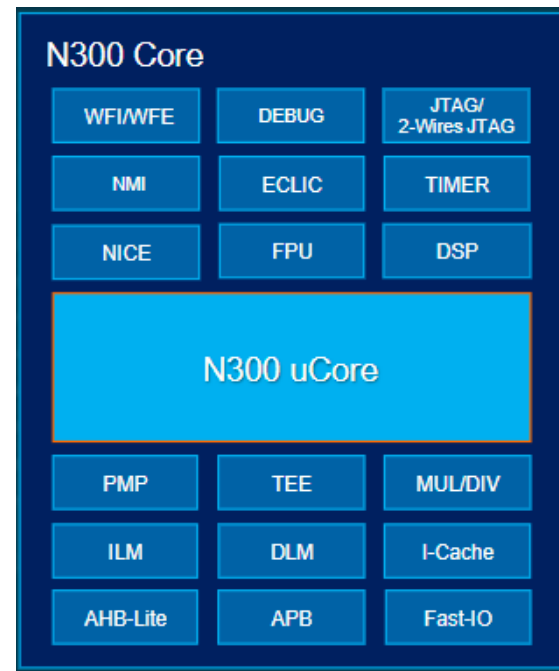
矢量

NN

# Nuclei Soc简述

SOC系列	RISC-V支持
N200	RV32I/E/M/A/C
N300	RV32I/E/M/A/C/F/D/P
N600	RV32I/M/A/C/F/D/P
NX600	RV64I/M/A/C/F/D/P
UX600(MMU)	RV64I/M/A/C/F/D/P
N900	RV32I/M/A/C/F/D/P/V
NX900	RV64I/M/A/C/F/D/P/V
UX900(MMU)	RV64I/M/A/C/F/D/P/V

- 工具链：[RISC-V GNU Toolchain](#)
- 软件环境—验证环境
  - [Nuclei SDK](#)
  - [Nuclei Linux SDK](#)
  - [RT-Thread Nuclei BSP](#)
- 文档来源：[Nuclei User Center](#)





# QEMU中NucLei Soc的添加

## ➤CPU模拟

- CPU Model
- 指令与CSR

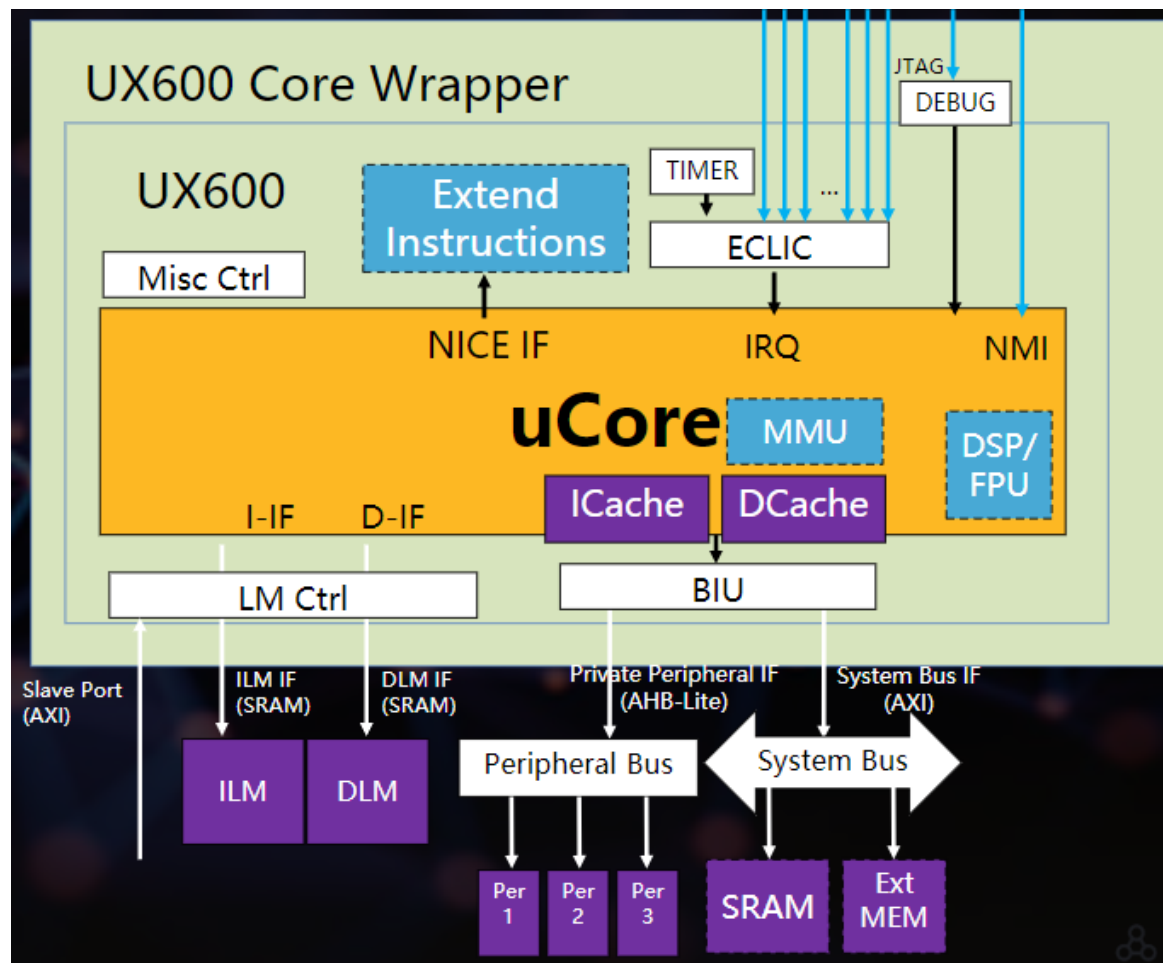
## ➤内存模拟

## ➤中断模拟

- ECLIC

## ➤外设模拟

- rcu
- systimer
- uart(示例)



UX600系统框架图示例

# QEMU Object Model (QOM)

## Features:

- System for dynamically registering types
- Support for single-inheritance of types
- Multiple inheritance of stateless interfaces

来自: [QOM DOC](#)

注册:

type_init(function)
module_init(function, type)
register_module_init(function, type);
QTAILQ_INSERT_TAIL(l, e, node);

```
typedef struct MyDevice
{
    DeviceState parent;
    int reg0, reg1, reg2;
} MyDevice;
```

```
typedef struct MyDeviceClass
{
    DeviceClass parent;
    void (*init) (MyDevice *obj);
} MyDeviceClass;
```

```
static void my_device_class_init(ObjectClass *oc, void *data)
{
}

static void my_device_init(Object *obj)
{
}
```

```
#define TYPE_MY_DEVICE "my-device"
static const TypeInfo my_device_info = {
    .parent = TYPE_DEVICE,
    .name = TYPE_MY_DEVICE,
    .instance_size = sizeof(MyDevice),
    .instance_init = my_device_init,
    .instance_finalize = my_device_finalize,
    .class_size = sizeof(MyDeviceClass),
    .class_init = my_device_class_init,
};
```

```
static void my_device_register_types(void)
{
    type_register_static(&my_device_info);
}
type_init(my_device_register_types)
```

```
DEFINE_TYPES(device_types_info)
```



# CPU虚拟化

Task 1 CPU Name 建立
Supported Cores: n201 n201e n203 n203e n205 n205e n305 n307 n307fd n600 n600f n600fd nx600 nx600f nx600fd ux600 ux600f ux600fd

Task 2 ADD Instructions
Instructions Extension: V-Extension for 1.0.0 P-Extension NICE(Nuclei Instruction Co- Unit Extension)

Task 3 ADD CSRs
CSR地址 读写属性 名称 0x307 MRW mttv 0x7eb MRW pushmsubm 0x7ec MRW mttv2 0x7ed MRW jalmnxti 0x7ee MRW pushmcause 0x7ef MRW pushmepc .....

# CPU模型

```
typedef struct RISCVCPU {  
    /*< private >*/  
    CPUState parent_obj;  
    /*< public >*/  
    CPUNegativeOffsetState neg;  
    CPURISCVState env;  
  
    /* Configuration Settings */  
    struct {  
    } cfg;  
} RISCVCPU;
```

```
typedef struct RISCVCPUClass {  
    /*< private >*/  
    CPUClass parent_class;  
    /*< public >*/  
    DeviceRealize parent_realize;  
    DeviceReset parent_reset;  
} RISCVCPUClass;
```

```
#define DEFINE_CPU(type_name, initfn) \  
    { \  
        .name = type_name, \  
        .parent = TYPE_RISCV_CPU, \  
        .instance_init = initfn \  
    }  
  
static const TypeInfo riscv_cpu_type_infos[] = {  
    {  
        .name = TYPE_RISCV_CPU,  
        .parent = TYPE_CPU,  
        .instance_size = sizeof(RISCVCPU),  
        .instance_init = riscv_cpu_init,  
        .abstract = true,  
        .class_size = sizeof(RISCVCPUClass),  
        .class_init = riscv_cpu_class_init,  
    },  
    DEFINE_CPU(TYPE_RISCV_CPU_ANY, riscv_any_cpu_init),  
#if defined(TARGET_RISCV32)  
    DEFINE_CPU(TYPE_RISCV_CPU_BASE32, riscv_base32_cpu_init),  
    DEFINE_CPU(TYPE_RISCV_CPU_SIFIVE_E31, rv32imacu_nommu_cpu_init),  
    .....,  
#elif defined(TARGET_RISCV64)  
    DEFINE_CPU(TYPE_RISCV_CPU_BASE64, riscv_base64_cpu_init),  
    DEFINE_CPU(TYPE_RISCV_CPU_SIFIVE_U54, rv64gcsu_priv1_10_0_cpu_init),  
    .....,  
#endif  
};  
  
DEFINE_TYPES(riscv_cpu_type_infos)
```

# CPU模型

```
typedef struct RISCVCPU {  
    /*< private >*/  
    CPUState parent_obj;  
    /*< public >*/  
    CPUNegativeOffsetState neg;  
    CPURISCVState env;  
  
    /* Configuration Settings */  
    struct {  
    } cfg;  
} RISCVCPU;
```

```
typedef struct RISCVCPUClass {  
    /*< private >*/  
    CPUClass parent_class;  
    /*< public >*/  
    DeviceRealize parent_realize;  
    DeviceReset parent_reset;  
} RISCVCPUClass;
```

```
#define TYPE_RISCV_CPU_BASE32      RISCV_CPU_TYPE_NAME("rv32")  
#define TYPE_RISCV_CPU_BASE64      RISCV_CPU_TYPE_NAME("rv64")  
  
static void rv32imafcu_nommu_cpu_init(Object *obj)  
{  
    CPURISCVState *env = &RISCV_CPU(obj)->env;  
    set_misa(env, RV32 | RVI | RVM | RVA | RVF | RVC | RVU);  
    set_priv_version(env, PRIV_VERSION_1_10_0);  
    set_resetvec(env, DEFAULT_RSTVEC);  
    set_feature(env, RISCV_FEATURE_PMP);  
}  
  
static void rv64gcsu_priv1_10_0_cpu_init(Object *obj)  
{  
    CPURISCVState *env = &RISCV_CPU(obj)->env;  
    set_misa(env, RV64 | RVI | RVM | RVA | RVF | RVD | RVC | RVS | RVU);  
    set_priv_version(env, PRIV_VERSION_1_10_0);  
    set_resetvec(env, DEFAULT_RSTVEC);  
    set_feature(env, RISCV_FEATURE_MMU);  
    set_feature(env, RISCV_FEATURE_PMP);  
}  
  
static void set_misa(CPURISCVState *env, target_ulong misa)  
{  
    env->misa_mask = env->misa = misa;  
}  
static void set_feature(CPURISCVState *env, int feature)  
{  
    env->features |= (1ULL << feature);  
}
```

# CPU模型

```
typedef struct RISCVCPU {  
    /*< private >*/  
    CPUState parent_obj;  
    /*< public >*/  
    CPUNegativeOffsetState neg;  
    CPURISCVState env;  
  
    /* Configuration Settings */  
    struct {  
    } cfg;  
} RISCVCPU;
```

```
typedef struct RISCVCPUClass {  
    /*< private >*/  
    CPUClass parent_class;  
    /*< public >*/  
    DeviceRealize parent_realize;  
    DeviceReset parent_reset;  
} RISCVCPUClass;
```

```
struct CPURISCVState {  
    target_ulong gpr[32];  
    uint64_t fpr[32]; /* assume both F and D extensions */  
    target_ulong pc;  
  
    target_ulong priv_ver;  
    target_ulong misa;  
    target_ulong misa_mask;  
    .....  
    uint32_t features;  
  
    target_ulong mstatus;  
    target_ulong mip;  
  
#ifdef TARGET_RISCV32  
    target_ulong mstatush;  
#endif  
    .....  
    target_ulong mie;  
    target_ulong mideleg;  
    .....  
    target_ulong sptbr; /* until: priv-1.9.1 */  
    target_ulong satp; /* since: priv-1.10.0 */  
    .....  
    /* Fields from here on are preserved across CPU reset. */  
    QEMUTimer *timer; /* Internal timer */  
};
```

# CPU模型

```
typedef struct RISCVCPU {  
    /*< private >*/  
    CPUState parent_obj;  
    /*< public >*/  
    CPUNegativeOffsetState neg;  
    CPURISCVState env;  
  
    /* Configuration Settings */  
    struct {  
    } cfg;  
} RISCVCPU;
```

```
typedef struct RISCVCPUClass {  
    /*< private >*/  
    CPUClass parent_class;  
    /*< public >*/  
    DeviceRealize parent_realize;  
    DeviceReset parent_reset;  
} RISCVCPUClass;
```

```
typedef struct CPUClass {  
    /*< private >*/  
    DeviceClass parent_class;  
    /*< public >*/  
  
    ObjectClass *(*class_by_name)(const char *cpu_model);  
    void (*parse_features)(const char *typename, char *str, Error **errp);  
  
    int reset_dump_flags;  
    bool (*has_work)(CPUState *cpu);  
    void (*do_interrupt)(CPUState *cpu);  
    .....  
    const VMStateDescription *vmsd;  
    const char *gdb_core_xml_file;  
    gchar * (*gdb_arch_name)(CPUState *cpu);  
    const char * (*gdb_get_dynamic_xml)(CPUState *cpu, const char *xmlname);  
    void (*cpu_exec_enter)(CPUState *cpu);  
    void (*cpu_exec_exit)(CPUState *cpu);  
    bool (*cpu_exec_interrupt)(CPUState *cpu, int interrupt_request);  
  
    void (*disas_set_info)(CPUState *cpu, disassemble_info *info);  
    vaddr (*adjust_watchpoint_address)(CPUState *cpu, vaddr addr, int len);  
    void (*tcg_initialize)(void);  
  
    /* Keep non-pointer data at the end to minimize holes. */  
    int gdb_num_core_regs;  
    bool gdb_stop_before_watchpoint;  
} CPUClass;
```

# CPU模型

以N307FD为示例:

```
#define TYPE_RISCV_CPU_NUCLEI_N307FD    RISCV_CPU_TYPE_NAME("nuclei-n307fd")

static const TypeInfo riscv_cpu_type_infos[] = {
    .....,
    DEFINE_CPU(TYPE_RISCV_CPU_NUCLEI_N307FD,    rv32imafdcu_nuclei_cpu_init),
    .....,
};

static void rv32imafdcu_nuclei_cpu_init(Object *obj)
{
    CPURISCVState *env = &RISCV_CPU(obj)->env;
    set_misa(env, RV32 | RVI | RVM | RVA | RVF | RVD | RVC | RVU);
    set_priv_version(env, PRIV_VERSION_1_10_0);
    set_resetvec(env, DEFAULT_RSTVEC);
    set_feature(env, RISCV_FEATURE_PMP);
}
```

运行结果:


```
wang@lelouch:~/workroom/qemu/plct-qemu/build$ ./riscv32-sofmmu/qemu-system-riscv32 -cpu ?
any
nuclei-n201
nuclei-n201e
nuclei-n203
nuclei-n203e
nuclei-n205
nuclei-n205e
nuclei-n305
nuclei-n307
nuclei-n307fd
nuclei-n600
nuclei-n600fd
```

```
wang@lelouch:~/workroom/qemu/plct-qemu/build64$ riscv64-sofmmu/qemu-system-riscv64 -cpu ?
any
nuclei-nx600
nuclei-nx600fd
rv64
```

# RISC-V CPU指令扩展

## Tiny Code Generator (TCG)

将源处理器机器代码转换为虚拟机运行所需的机器代码块(如x86机器代码块)



主要调用栈:

```
cpu_loop(CPURISCVState *env)
cpu_exec(CPUState *cpu)
TranslationBlock *tb_find(.....)
TranslationBlock *tb_gen_code(.....)
void gen_intermediate_code(.....)
translator_loop(const TranslatorOps *ops, ..... )
tcg_gen_code(TCGContext *s, TranslationBlock *tb)
```

target/arch/translate.c

```
struct TranslatorOps
{
    void (*init_disas_context)
    void (*tb_start)
    void (*insn_start)
    bool (*breakpoint_check)
    void (*translate_insn)
    void (*tb_stop)
    void (*disas_log)
```



```
void decode_opc()
decode_insn16(ctx, opcode)
decode_insn32(ctx, opcode32)
```

} decodeTree



# RISC-V CPU指令扩展

## ADDI Instructions


ADDI adds the **sign-extended** 12-bit immediate to register rs1. Arithmetic overflow is ignored and the result is simply the low XLEN bits of the result.

ADDI rd, rs1, 0 is used to implement the **MV rd, rs1** assembler pseudoinstruction.

来自: [Unprivileged ISA](#)

```
addi rd, rs1, immediate
x[rd] = x[rs1] + sext(immediate)
```

```
typedef arg_i arg_addi;
static bool trans_addi(DisasContext *ctx, arg_addi *a);
```



target/riscv/insn32.decode

target/riscv/decode\_insn32.inc.c

# Fields:

%rs1 15:5

%rd 7:5

# immediates:

%imm\_i 20:s12

# Argument sets:

&i imm rs1 rd

# Formats 32:

@i ..... &i imm=%imm\_i %rs1 %rd

# \*\*\* RV32I Base Instruction Set \*\*\*

addi ..... 000 ..... 0010011 @i

```
static void decode_insn32_extract_i(D
isasContext *ctx, arg_i *a, uint32_t
insn)
{
    a->imm = sext32(insn, 20, 12);
    a->rs1 = extract32(insn, 15, 5);
    a->rd = extract32(insn, 7, 5);
}
```

```
typedef struct
{
    int imm;
    int rd;
    int rs1;
} arg_i;
```

# RISC-V CPU指令扩展

target/riscv/decode\_insn32.inc.c

```
static bool decode_insn32(DisasContext *ctx, uint32_t insn)
{
    switch (insn & 0x0000007f) {
        .....
        case 0x00000013:
            /* ..... .0010011 */
            switch ((insn >> 12) & 0x7) {
                case 0x0:
                    /* ..... .000.... .0010011 */
                    decode_insn32_extract_i(ctx, &u.f_i, insn);
                    if (trans_addi(ctx, &u.f_i)) return true;
                    return false;
                .....
            }
        .....
    }
}
```

target/riscv/insn\_trans/trans\_rvi.inc.c

```
static bool trans_addi(DisasContext *ctx, arg_addi *a)
{
    return gen_arith_imm_fn(ctx, a, &tcg_gen_addi_tl);
}

#define tcg_gen_addi_tl tcg_gen_addi_i32

void tcg_gen_addi_i32(TCGv_i32 ret, TCGv_i32 arg1, int32_t arg2)
{
    /* some cases can be optimized here */
    if (arg2 == 0) {
        tcg_gen_mov_i32(ret, arg1);
    } else {
        TCGv_i32 t0 = tcg_const_i32(arg2);
        tcg_gen_add_i32(ret, arg1, t0);
        tcg_temp_free_i32(t0);
    }
}
```

参考:  
[Backend Ops](#)  
[Frontend Ops](#)

# Nuclei Soc CSR扩展

N级别处理器内核自定义CSR部分:

0x7d9	MRW	msaveepc2	自定义寄存器用于保存第二级嵌套 NMI 或异常的 mepc 注意: 此寄存器只有配置了两级异常嵌套恢复才会存在
0x7da	MRW	msavecause2	自定义寄存器用于保存第二级嵌套 NMI 或异常的 mcause 注意: 此寄存器只有配置了两级异常嵌套恢复才会存在
0x7db	MRW	msavedcause1	自定义寄存器用于保存第一级嵌套异常的 mdcouse 注意: 此寄存器只有配置了两级异常嵌套恢复才会存在
0x7dc	MRW	msavedcause2	自定义寄存器用于保存第二级嵌套异常的 mdcouse 注意: 此寄存器只有配置了两级异常嵌套恢复才会存在
0x7eb	MRW	pushmsubm	自定义寄存器用于将 msubm 的值存入堆栈地址空间
0x7ec	MRW	mtvt2	自定义寄存器用于设定非向量中断处理模式的中断入口地址
0x7ed	MRW	jalmnxti	自定义寄存器用于使能 ECLIC 中断, 该寄存器的读操作能处理下一个中断同时返回下一个中断 Handler 的入口地址, 并跳转至此地址。

来自:《Nuclei\_N级别指令架构手册》

```
* csrr    <->  riscv_csrrw(env, csrno, ret_value, 0, 0);
* csrrw   <->  riscv_csrrw(env, csrno, ret_value, value, -1);
* csrrs   <->  riscv_csrrw(env, csrno, ret_value, -1, value);
* csrrc   <->  riscv_csrrw(env, csrno, ret_value, 0, value);
```

target/riscv/insn32.decode

```
# Fields:
%rs1    15:5
%rd     7:5
%csr    20:12
# Formats 32:
@csr    .....  ....  ...  .....          %csr  %rs1 %rd

# *** RV32I Base Instruction Set ***
csrrw    .....  .... 001 ..... 1110011 @csr
csrrs    .....  .... 010 ..... 1110011 @csr
csrrc    .....  .... 011 ..... 1110011 @csr
csrrwi   .....  .... 101 ..... 1110011 @csr
csrrsi   .....  .... 110 ..... 1110011 @csr
csrrci   .....  .... 111 ..... 1110011 @csr
```



```
static bool trans_csrrw(DisasContext *ctx, arg_csrrw *a)
static bool trans_csrrs(DisasContext *ctx, arg_csrrs *a)
static bool trans_csrrc(DisasContext *ctx, arg_csrrc *a)
static bool trans_csrrwi(DisasContext *ctx, arg_csrrwi *a)
static bool trans_csrrsi(DisasContext *ctx, arg_csrrsi *a)
static bool trans_csrrci(DisasContext *ctx, arg_csrrci *a)
```

```
gen_helper_csrrw
gen_helper_csrrs
gen_helper_csrrc
```



```
DEF_HELPER_3(csrrw, tl, env, tl, tl)
DEF_HELPER_4(csrrs, tl, env, tl, tl, tl)
DEF_HELPER_4(csrrc, tl, env, tl, tl, tl)
```

# Nuclei Soc CSR扩展

target/riscv/cpu\_bits.h

```
#define CSR_PUSHMSUBM    0x07eb  
#define CSR_MTVT2        0x07ec  
#define CSR_JALMNXTI     0x07ed  
#define CSR_PUSHMCAUSE   0x07ee  
#define CSR_PUSHMEPC     0x07ef
```



target/riscv/csr.c

```
/* Control and Status Register function table */  
static riscv_csr_operations csr_ops[CSR_TABLE_SIZE] = {  
    /* User Floating-Point CSRs */  
    [CSR_FFLAGS] = { fs, read_fflags, write_fflags },  
    [CSR_FRM] = { fs, read_frm, write_frm },  
    [CSR_FCSR] = { fs, read_fcsr, write_fcsr },  
  
    /* User Timers and Counters */  
    [CSR_CYCLE] = { ctr, read_instret },  
    [CSR_INSTRET] = { ctr, read_instret },  
    .....  
};
```

```
struct riscv_csr_operations
```

```
riscv_csr_predicate_fn predicate
```

```
riscv_csr_read_fn read
```

```
riscv_csr_write_fn write
```

```
riscv_csr_op_fn op
```

## CSR PUSHMCAUSE:

处理器定义了通过 pushmcause 寄存器 csrrwi 操作实现的 CSR 指令，存储 mcause 的值到堆栈指针作为基地址的memory 空间  
以如下指令为例介绍此CSR指令：

csrrwi x0, PUSHMCAUSE, 1

该指令的操作是将mcause寄存器的值存到SP（堆栈指针）+1\*4的地址。

```
static int rmw_pushmcause(CPURISCVState *env, int csrno, target_ulong *ret_value,  
                          target_ulong new_value, target_ulong write_mask)  
{  
    uint64_t notify_addr = new_value * 4 + env->gpr[2];  
    .....  
    cpu_physical_memory_rw(notify_addr, &env->mcause, 4, 1);  
    .....  
    return 0;  
}
```

# 内存虚拟化

The memory API models the **memory** and **I/O buses** and **controllers** of a QEMU machine.

- ordinary RAM
- memory-mapped I/O (MMIO)
- memory controllers that can dynamically reroute physical memory regions to different destinations

来自: [memory API](#)

Types of regions	initialize
RAM	memory_region_init_ram()
MMIO	memory_region_init_io()
ROM	memory_region_init_rom().
ROM device	memory_region_init_rom_device()
IOMMU region	memory_region_init_iommu()
container	memory_region_init()
alias	memory_region_init_alias()
reservation region	memory_region_init_io()

## MemoryRegion

```
QTAILQ_HEAD(, MemoryRegion) subregions  
  
const MemoryRegionOps *ops  
  
MemoryRegion *alias  
  
hwaddr alias_offset
```

# 内存虚拟化

MemoryRegion ROOT — system\_memory:

```
MemoryRegion *get_system_memory(void)
```

RAM:

```
void memory_region_init_ram(MemoryRegion *mr,  
                             struct Object *owner,  
                             const char *name,  
                             uint64_t size,  
                             Error **errp)
```

ROM:

```
void memory_region_init_rom(MemoryRegion *mr,  
                             struct Object *owner,  
                             const char *name,  
                             uint64_t size,  
                             Error **errp)
```

分配/挂载:

```
void memory_region_add_subregion(MemoryRegion *mr,  
                                 hwaddr offset,  
                                 MemoryRegion *subregion)
```

IO:

```
void memory_region_init_io(MemoryRegion *mr,  
                             Object *owner,  
                             const MemoryRegionOps *ops,  
                             void *opaque,  
                             const char *name,  
                             uint64_t size)  
void sysbus_init_mmio(SysBusDevice *dev, MemoryRegion *memory)  
void sysbus_mmio_map(SysBusDevice *dev, int n, hwaddr addr)
```

```
/*  
 * Memory region callbacks  
 */  
struct MemoryRegionOps {  
    /* Read from the memory region. @addr is relative to  
    @mr; @size is  
    * in bytes. */  
    uint64_t (*read)(void *opaque,  
                      hwaddr addr,  
                      unsigned size);  
    /* Write to the memory region. @addr is relative to @  
    mr; @size is  
    * in bytes. */  
    void (*write)(void *opaque,  
                  hwaddr addr,  
                  uint64_t data,  
                  unsigned size);  
    .....  
}
```

# 内存虚拟化

Table 5-1 Address Allocation of SoC

	Component	Address Spaces	Description
Core Private Peripherals	TIMER	0x0200_0000 ~ 0x0200_0FFF	TIMER Unit address space.
	ECLIC	0x0C00_0000 ~ 0x0C00_FFFF	ECLIC Unit address space.
	DEBUG	0x0000_0000 ~ 0x0000_0FFF	DEBUG Unit address space.
Memory Resource	ILM	0x8000_0000 ~	ILM address space.
	DLM	0x9000_0000 ~	DLM address space.
	ROM	0x0000_1000 ~ 0x0000_1FFF	Internal ROM.
	Off-Chip QSPIo Flash Read	0x2000_0000 ~ 0x3FFF_FFFF	QSPIo with XiP mode read-only address space.
Peripherals	GPIO	0x1001_2000 ~ 0x1001_2FFF	GPIO Unit address space.
	UART0	0x1001_3000 ~ 0x1001_3FFF	First UART address space.
	QSPIo	0x1001_4000 ~ 0x1001_4FFF	First QSPI address space.
	PWM0	0x1001_5000 ~ 0x1001_5FFF	First PWM address space.
	UART1	0x1002_3000 ~ 0x1002_3FFF	Second UART address space.
	QSPI1	0x1002_4000 ~ 0x1002_4FFF	Second QSPI address space.
	PWM1	0x1002_5000 ~ 0x1002_5FFF	Second PWM address space.
	QSPI2	0x1003_4000 ~ 0x1003_4FFF	Third QSPI address space.
	PWM2	0x1003_5000 ~ 0x1003_5FFF	Third PWM address space.
	I2C Master	0x1004_2000 ~ 0x1004_2FFF	I2C Master address space.
Default slave	The other space is write-ignored and read-as zero.		

```
static const struct MemmapEntry {
    hwaddr base;
    hwaddr size;
} nuclei_memmap[] = {
    [NUCLEI_DEBUG] = { 0x0, 0x1000 },
    [NUCLEI_ROM] = { 0x1000, 0x1000 },
    [NUCLEI_TIMER] = { 0x2000000, 0x1000 },
    [NUCLEI_ECLIC] = { 0xc000000, 0x10000 },
    [NUCLEI_GPIO] = { 0x10012000, 0x1000 },
    [NUCLEI_UART0] = { 0x10013000, 0x1000 },
    [NUCLEI_QSPI0] = { 0x10014000, 0x1000 },
    [NUCLEI_PWM0] = { 0x10015000, 0x1000 },
    [NUCLEI_UART1] = { 0x10023000, 0x1000 },
    [NUCLEI_QSPI1] = { 0x10024000, 0x1000 },
    [NUCLEI_PWM1] = { 0x10025000, 0x1000 },
    [NUCLEI_QSPI2] = { 0x10034000, 0x1000 },
    [NUCLEI_PWM2] = { 0x10035000, 0x1000 },
    [NUCLEI_XIP] = { 0x20000000, 0x10000000 },
    [NUCLEI_ILM] = { 0x80000000, 0x20000 },
    [NUCLEI_DLM] = { 0x90000000, 0x20000 },
};
```



# 内存虚拟化

Table 5-1 Address Allocation of SoC

	Component	Address Spaces	Description
Core Private Peripherals	TIMER	0x0200_0000 ~ 0x0200_0FFF	TIMER Unit address space.
	ECLIC	0x0C00_0000 ~ 0x0C00_FFFF	ECLIC Unit address space.
	DEBUG	0x0000_0000 ~ 0x0000_0FFF	DEBUG Unit address space.
Memory Resource	ILM	0x8000_0000 ~	ILM address space.
	DLM	0x9000_0000 ~	DLM address space.
	ROM	0x0000_1000 ~ 0x0000_1FFF	Internal ROM.
	Off-Chip QSPIo Flash Read	0x2000_0000 ~ 0x3FFF_FFFF	QSPIo with XiP mode read-only address space.
Peripherals	GPIO	0x1001_2000 ~ 0x1001_2FFF	GPIO Unit address space.
	UARTo	0x1001_3000 ~ 0x1001_3FFF	First UART address space.
	QSPIo	0x1001_4000 ~ 0x1001_4FFF	First QSPI address space.
	PWMo	0x1001_5000 ~ 0x1001_5FFF	First PWM address space.
	UART1	0x1002_3000 ~ 0x1002_3FFF	Second UART address space.
	QSPI1	0x1002_4000 ~ 0x1002_4FFF	Second QSPI address space.
	PWM1	0x1002_5000 ~ 0x1002_5FFF	Second PWM address space.
	QSPI2	0x1003_4000 ~ 0x1003_4FFF	Third QSPI address space.
	PWM2	0x1003_5000 ~ 0x1003_5FFF	Third PWM address space.
	I2C Master	0x1004_2000 ~ 0x1004_2FFF	I2C Master address space.
Default slave	The other space is write-ignored and read-as zero.		

ILM DLM示例:

```
MemoryRegion *sys_mem = get_system_memory();
```

```
memory_region_init_ram(&s->soc.ilm, NULL, "riscv.nuclei.ram.ilm",  
                        memmap[NUCLEI_ILM].size, &error_fatal);  
memory_region_add_subregion(sys_mem,  
                             memmap[NUCLEI_ILM].base, &s->soc.ilm);
```

```
memory_region_init_ram(&s->soc.dlm, NULL, "riscv.nuclei.ram.dlm",  
                        memmap[NUCLEI_DLM].size, &error_fatal);  
memory_region_add_subregion(sys_mem,  
                             memmap[NUCLEI_DLM].base, &s->soc.dlm);
```

```
/* Mask ROM */
```

```
memory_region_init_rom(&s->internal_rom, OBJECT(dev), "riscv.nuclei.  
irom", memmap[NUCLEI_ROM].size, &error_fatal);  
memory_region_add_subregion(sys_mem,  
                             memmap[NUCLEI_ROM].base, &s->internal_rom);
```

```
hw/riscv/sifive_gpio.c
```

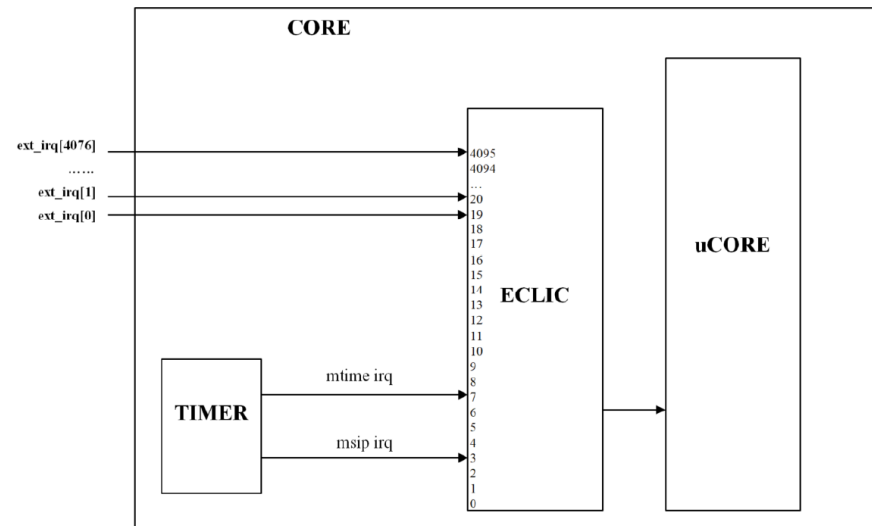
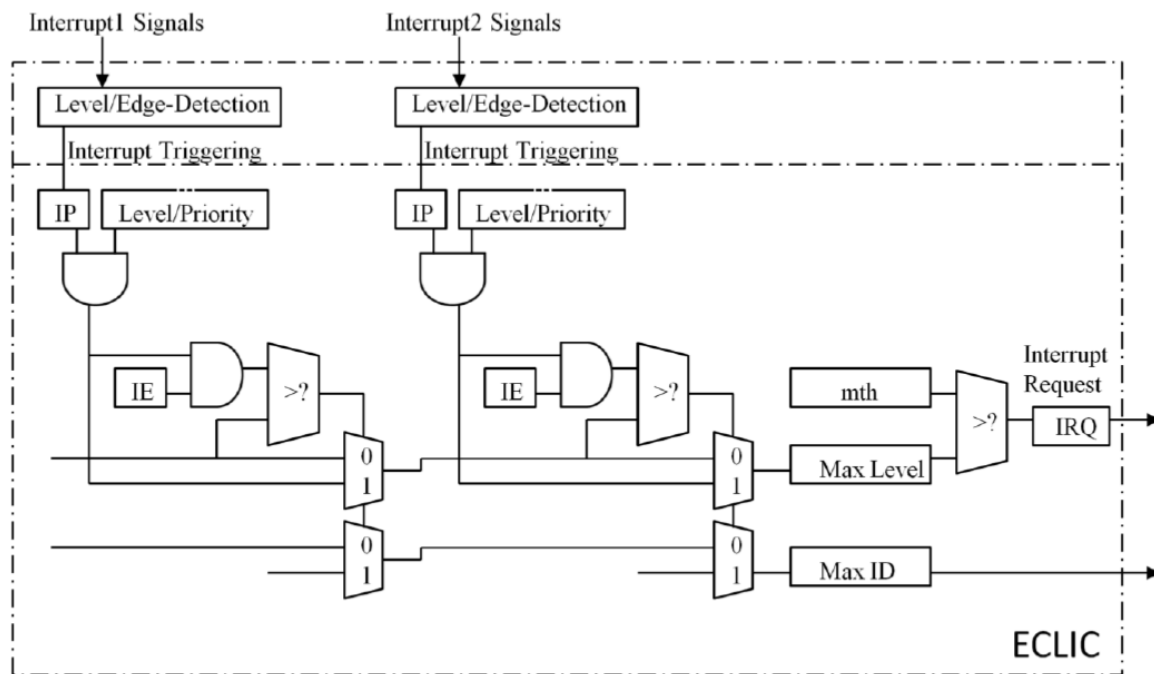
```
memory_region_init_io(&s->mmio, obj, &gpio_ops, s,  
                      TYPE_SIFIVE_GPIO, SIFIVE_GPIO_SIZE);  
sysbus_init_mmio(SYS_BUS_DEVICE(obj), &s->mmio);
```

```
hw/riscv/nuclei_hbrid.c
```

```
object_property_set_bool(OBJECT(&s->gpio), true, "realized", &err);  
sysbus_mmio_map(SYS_BUS_DEVICE(&s->gpio), 0, memmap[NUCLEI_GPIO].base);
```

# 中断虚拟化

## ECLIC: Enhanced Core Local Interrupt Controller



- ECLIC 只服务于一个处理器内核，为该处理器内核私有
- 兼容CLIC
- 支持4096个中断源
- 支持向量和非向量处理
- 支持中断响应、嵌套、咬尾
- .....

# 中断虚拟化

表 6-6 ECLIC 寄存器的单元内地址偏移量

	属性	名称	宽度
0x0000	可读可写	clccfg	8 位
0x0004	只读, 写忽略	clicinfo	32 位
0x000b	可读可写	mth	8 位
0x1000+4*i	可读可写	clicintip[i]	8 位
0x1001+4*i	可读可写	clicintie[i]	8 位
0x1002+4*i	可读可写	clicintattr[i]	8 位
0x1003+4*i	可读可写	clicintctl[i]	8 位

```
struct IRQState {  
    Object parent_obj;  
  
    qemu_irq_handler handler;  
    void *opaque;  
    int n;  
};  
typedef struct IRQState *qemu_irq;
```

```
static inline void cpu_interrupt(  
    CPUState *cpu, int mask)
```

```
void cpu_reset_interrupt(CPUState  
    *cpu, int mask)
```

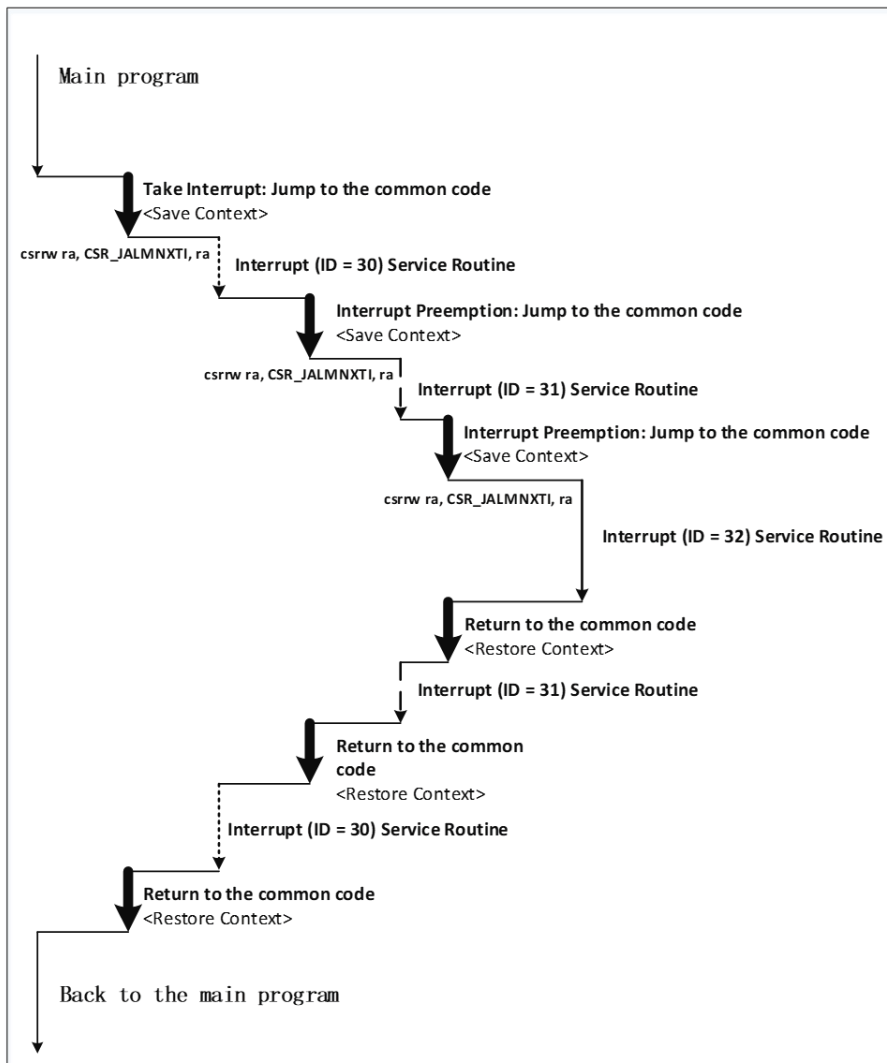
```
qemu_irq qemu_allocate_irq(qemu_irq_handler handler, void *opaque, int n)
```

```
void riscv_cpu_do_interrupt(CPUState *cs)
```

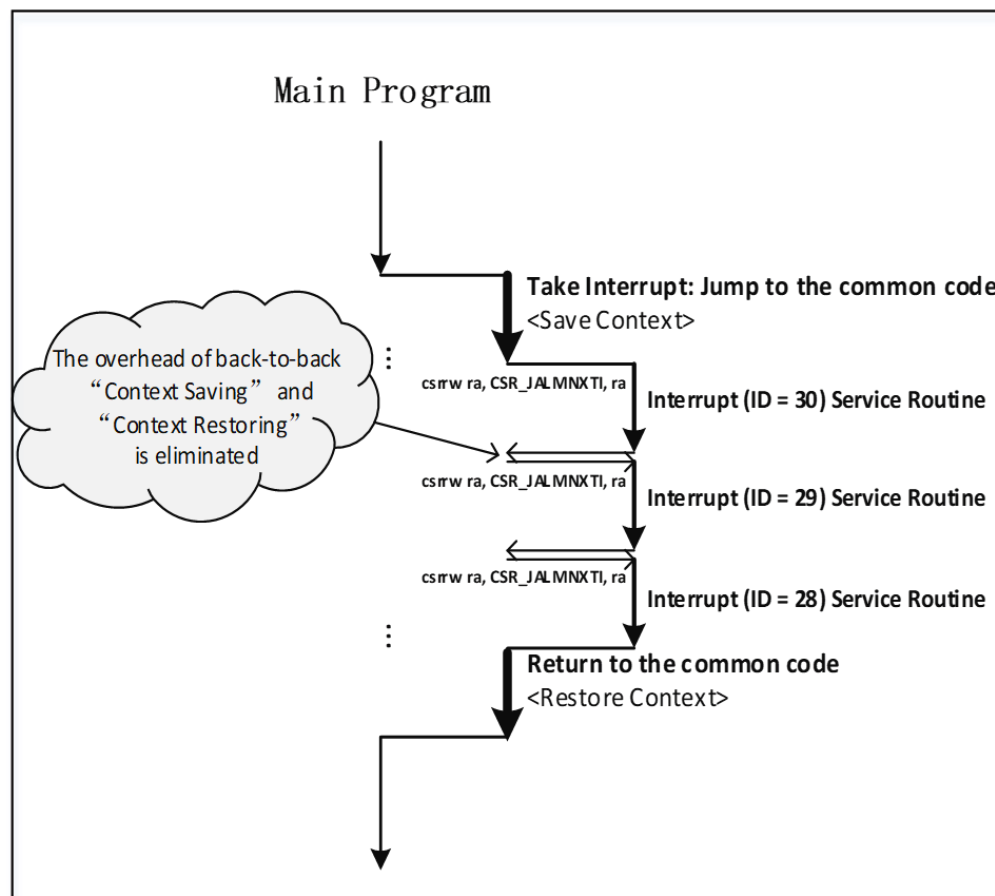
```
typedef struct NucLeiECLICState {  
    /*< private >*/  
    SysBusDevice parent_obj;  
  
    /*< public >*/  
    MemoryRegion mmio;  
    uint32_t num_sources;  
  
    /* config */  
    uint32_t sources_id;  
    uint8_t cliccfg;  
    uint32_t clicinfo;  
    uint8_t mth;  
    uint8_t *clicintip;  
    uint8_t *clicintie;  
    uint8_t *clicintattr;  
    uint8_t *clicintctl;  
    ECLICPendingInterrupt *clicintlist;  
    uint32_t aperture_size;  
  
    QLIST_HEAD(, ECLICPendingInterrupt) pending_list;  
    size_t active_count;  
  
    /* ECLIC IRQ handlers */  
    qemu_irq *irqs;  
}  
NucLeiECLICState;
```

```
void qemu_set_irq(qemu_irq irq, int level);
```

# 中断虚拟化(非向量化)



嵌套流程



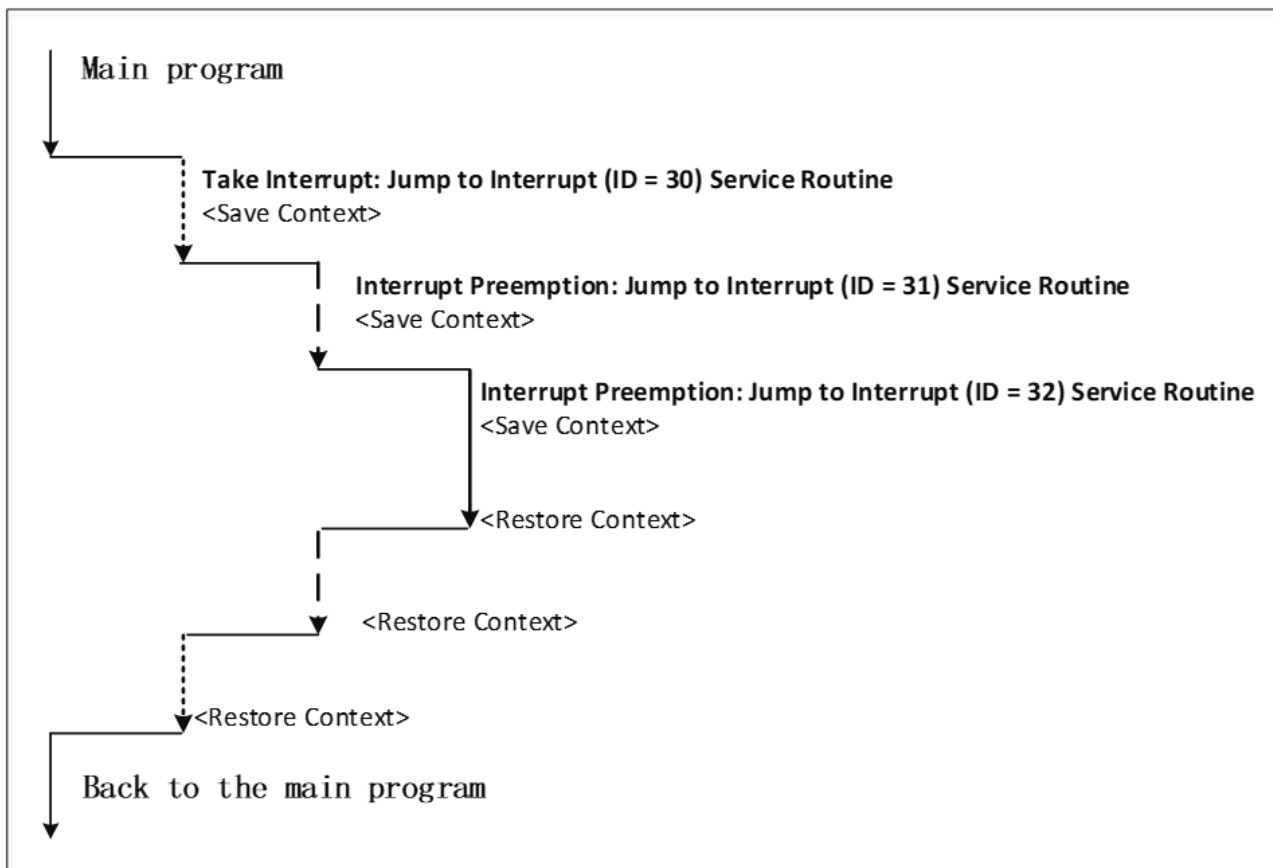
咬尾流程

irq\_entry入口函数:

```
sw    ra,0(sp)
.....
csrwi  0x7ee,11
csrwi  0x7ef,12
csrwi  0x7eb,13
csrrw  ra,0x7ed,ra
csrci  mstatus,8
csrw   0x7c4,t0
csrw   mepc,t0
.....
lw     ra,0(sp)
.....
mret
```

```
0x7c4 MRW msubm
0x7eb MRW pushmsubm
0x7ec MRW mtvt2
0x7ed MRW jalmnxti
0x7ee MRW pushmcause
0x7ef MRW pushmepc
```

# 中断虚拟化(向量化)



嵌套流程

## mstatus的MIE域

mstatus

寄存器中的 MIE 域表示机器模式下的全局中断使能:

当MIE 域的值 为 1 时, 表示中断的全局开关打开, 中断能够被正常响应;

当MIE 域的值 为 0 时, 表示全局关闭中断, 中断被屏蔽, 无法被响应。

## nuclei-sdk/demo\_eclic

```
__INTERUPT void eclic_msip_handler(void)
{
    static uint32_t int_sw_cnt = 0; /* software i
nterrupt counter */
    // save CSR context
    SAVE_IRQ_CSR_CONTEXT();
    SysTimer_ClearSWIRQ();
    .....
    // restore CSR context
    RESTORE_IRQ_CSR_CONTEXT();
}
```

### eclic\_msip\_handler

```
8000115a <eclic_msip_handler>:
8000115c: df86                                sw ra,252(sp)
)
8000117e: dd7e                                sw t6,184(sp)
80001180: b502                                fsd ft0,168(sp)
800011a6: a87e                                fsd ft11,16(sp)
800011aa: 342027f3                           csrr a5,mcause
800011b6: 341027f3                           csrr a5,mepc
800011c2: 7c4027f3                           csrr a5,0x7c4
800011ce: 40a1                                li ra,8
800011d0: 3000a073                           csrs mstatus,ra
80001224: 3004b073                           csrr mstatus,s1
8000122e: 7c449073                           csrw 0x7c4,s1
80001236: 34149073                           csrr mepc,s1
8000123e: 34249073                           csrr mcause,s1
80001244: 50fe                                lw ra,252(sp)
80001266: 5fea                                lw t6,184(sp)
80001268: 302a                                fld ft0,168(sp)
8000128e: 2fc2                                fld ft11,16(sp)
80001292: 30200073                           mret
```

# 外设虚拟化—UART

寄存器名称	偏移地址	描述
UART_TXDATA	0x000	发送数据寄存器
UART_RXDATA	0x004	接收数据寄存器
UART_TXCTRL	0x008	发送控制寄存器
UART_RXCTRL	0x00C	接收控制寄存器
UART_IE	0x010	UART中断使能寄存器
UART_IP	0x014	UART中断等待标志寄存器
UART_DIV	0x018	波特率生成分频系数

```
static const MemoryRegionOps uart_ops = {  
    .read = uart_read,  
    .write = uart_write,  
    .endianness = DEVICE_NATIVE_ENDIAN,  
    .valid = {  
        .min_access_size = 4,  
        .max_access_size = 4  
    }  
};
```

## UART后端设置:

```
qemu_chr_fe_init(&s->chr, chr, &error_abort);  
qemu_chr_fe_set_handlers(&s->chr, uart_can_rx, uart_rx, uart_event,  
    uart_be_change, s, NULL, true);
```

## UART描述:

```
typedef struct NucLeiUARTState {  
    /*< private >*/  
    SysBusDevice parent_obj;  
  
    /*< public >*/  
    qemu_irq irq;  
    MemoryRegion mmio;  
    CharBackend chr;  
    uint8_t rx_fifo[8];  
    unsigned int rx_fifo_len;  
  
    uint32_t txdata;  
    uint32_t rxdata;  
    uint32_t txctrl;  
    uint32_t rxctrl;  
    uint32_t ie;  
    uint32_t ip;  
    uint32_t div;  
} NucLeiUARTState;
```

# Uart 实现

寄存器读操作:

```
static uint64_t
uart_read(void *opaque, hwaddr offset, unsigned int size)
{
    NucleiUARTState *s = opaque;
    uint64_t value = 0;
    uint8_t fifo_val;

    switch (offset)
    {
        case NUCLEI_UART_REG_TXDATA:
            return 0;
        case NUCLEI_UART_REG_RXDATA:
            if (s->rx_fifo_len) {
                fifo_val = s->rx_fifo[0];
                memmove(s->rx_fifo, s->rx_fifo + 1, s->rx_fifo_len - 1);
                s->rx_fifo_len--;
                qemu_chr_fe_accept_input(&s->chr);
                update_irq(s);
                return fifo_val;
            }
            return 0x80000000;
        case NUCLEI_UART_REG_TXCTRL:
            value = s->txctrl;
            break;
        .....
    }
}
```

寄存器写操作:

```
static void
uart_write(void *opaque, hwaddr offset,
            uint64_t value, unsigned int size)
{
    NucleiUARTState *s = opaque;
    unsigned char ch = value;

    switch (offset)
    {
        case NUCLEI_UART_REG_TXDATA:
            qemu_chr_fe_write(&s->chr, &ch, 1);
            update_irq(s);
            break;
        case NUCLEI_UART_REG_TXCTRL:
            s->txctrl = value;
            break;
        .....
    }
}
```

UART中断:


```
static void update_irq(NucleiUARTState *s)
{
    int cond = 0;
    s->txctrl |= 0x1;
    if (s->rx_fifo_len)
        s->rxctrl &= ~0x1;
    else
        s->rxctrl |= 0x1;

    if ((s->ie & NUCLEI_UART_IE_TXWM) ||
        (s->ie & NUCLEI_UART_IE_RXWM) && s->rx_fifo_len)) {
        cond = 1;
    }

    if (cond) {
        qemu_irq_raise(s->irq);
    } else {
        qemu_irq_lower(s->irq);
    }
}
```



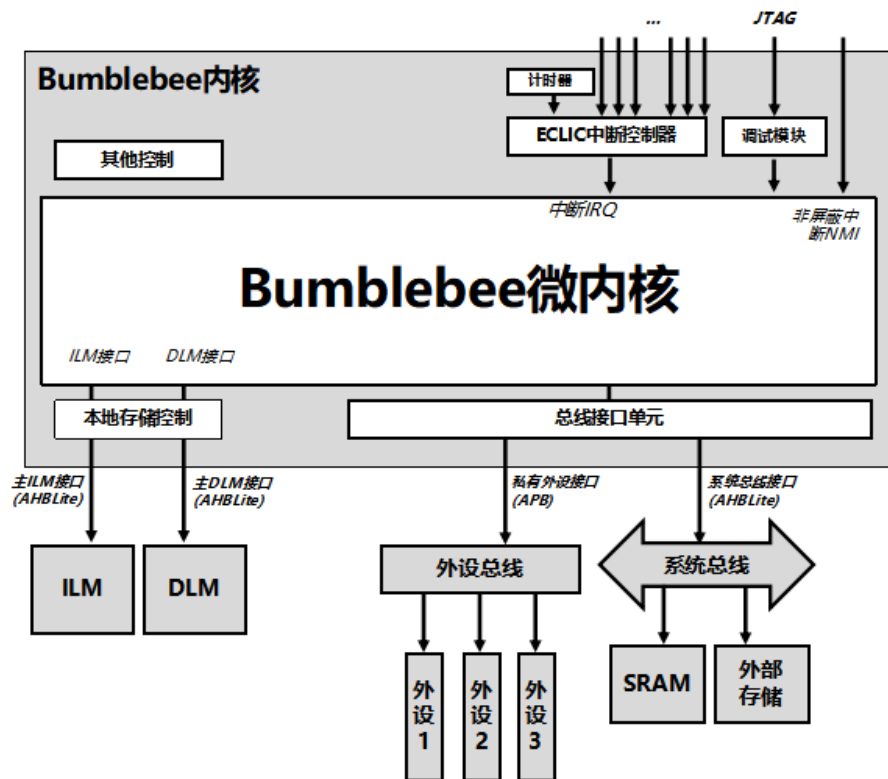
# GD32V103F Board



微控制器: GD32VF103VBT6 (32位RISC-V处理器)  
内核: 芯来科技Bumblebee内核 (RV32IMAC)  
主频: 108MHz  
内存: 内置128KB Flash、32KB SRAM  
工作电压: 2.6~3.6V  
外设资源: Timer (高级16位定时器\*1, 通用16位定时器\*4)  
U(S)ART\*5、I2C\*2、SPI\*3、CAN\*2、USBFS\*1、  
ADC\*2 (16路外部通道)、DAC\*2、EXMC\*1

[gd32v103f](#)

# GD32V103F Board



[gd32v103f](#)

微控制器: GD32VF103VBT6 (32位RISC-V处理器)

内核: 芯来科技Bumblebee内核 (RV32IMAC)

主频: 108MHz

内存: 内置128KB Flash、32KB SRAM

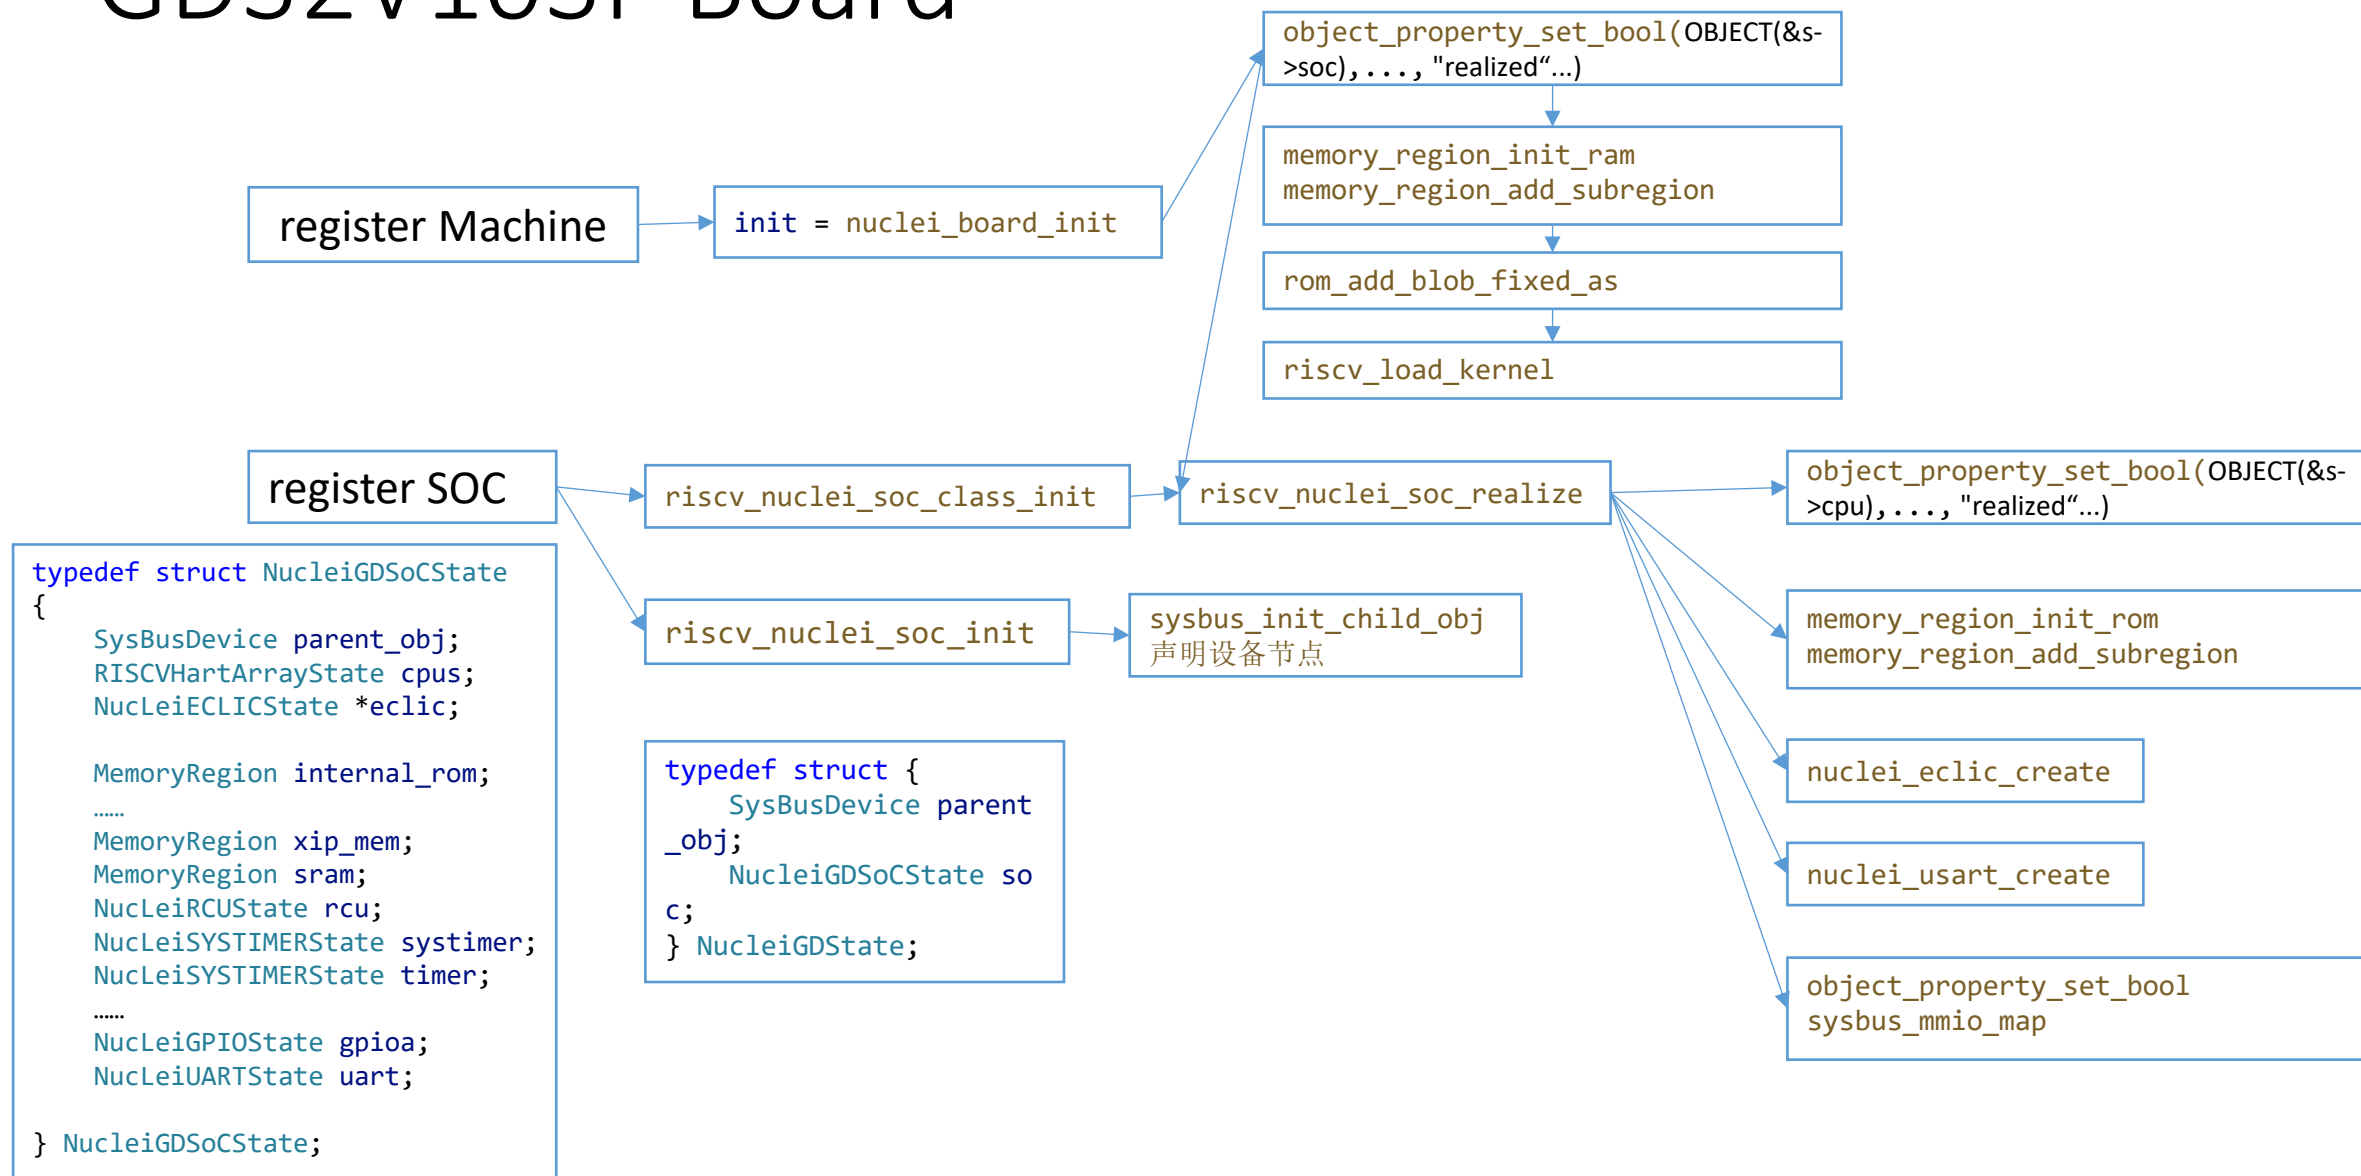
工作电压: 2.6~3.6V

外设资源: Timer (高级16位定时器\*1, 通用16位定时器\*4)

U(S)ART\*5、I2C\*2、SPI\*3、CAN\*2、USBFS\*1、

ADC\*2 (16路外部通道)、DAC\*2、EXMC\*1

# GD32V103F Board



# GD32V103F Board

```
qemu-system-riscv32 -nographic -machine gd32vf103_rvstar \  
-kernel ../os/rtthread/rt-thread/bsp/nuclei/gd32vf103_rvstar/rtthread.elf -nodefaults -serial stdio
```

```
wang@lelouch:~/workroom/qemu$ ./out/qemu-nuclei32/bin/qemu-system-riscv32 -nographic -machine gd32vf103_rvstar -kernel ../os/  
rtthread/rt-thread/bsp/nuclei/gd32vf103_rvstar/rtthread.elf -nodefaults -serial stdio  
initialize rti_board_start:0 done  
  
 \ | /  
- RT -   Thread Operating System  
 / | \   4.0.3 build Sep  9 2020  
2006 - 2020 Copyright by rt-thread team  
do components initialization.  
initialize rti_board_end:0 done  
initialize rt_work_sys_workqueue_init:0 done  
initialize rt_hw_pin_init:0 done  
initialize libc_system_init:0 done  
initialize finsh_system_init:0 done  
msh >ps  
thread  pri  status      sp      stack size max used left tick  error  
-----  
tshell   20  running 0x000000f8 0x00000800    21% 0x0000000a 000  
sys_work 23  suspend 0x00000098 0x00000800    07% 0x0000000a 000  
tidle0   31  ready  0x00000088 0x0000018c    34% 0x00000006 000  
timer     4  suspend 0x00000098 0x00000200    29% 0x0000000a 000
```

## GD32V103F Board

Final values of the variables used in the benchmark:

**dhystone**

```
Int_Glob: 5
  should be: 5
Bool_Glob: 1
  should be: 1
Ch_1_Glob: A
  should be: A
Ch_2_Glob: B
  should be: B
Arr_1_Glob[8]: 7
  should be: 7
Arr_2_Glob[8][7]: 5000010
  should be: Number_Of_Runs + 10
Ptr_Glob->
  Ptr_Comp: 536883328
  should be: (implementation-dependent)
Discr: 0
  should be: 0
Enum_Comp: 2
  should be: 2
Int_Comp: 17
  should be: 17
Str_Comp: DHRYSTONE PROGRAM, SOME STRING
  should be: DHRYSTONE PROGRAM, SOME STRING
Next_Ptr_Glob->
  Ptr_Comp: 536883328
  should be: (implementation-dependent), same as above
Discr: 0
  should be: 0
Enum_Comp: 1
  should be: 1
Int_Comp: 18
  should be: 18
Str_Comp: DHRYSTONE PROGRAM, SOME STRING
  should be: DHRYSTONE PROGRAM, SOME STRING
Int_1_Loc: 5
  should be: 5
Int_2_Loc: 13
  should be: 13
Int_3_Loc: 7
  should be: 7
Enum_Loc: 1
  should be: 1
Str_1_Loc: DHRYSTONE PROGRAM, 1'ST STRING
  should be: DHRYSTONE PROGRAM, 1'ST STRING
Str_2_Loc: DHRYSTONE PROGRAM, 2'ND STRING
  should be: DHRYSTONE PROGRAM, 2'ND STRING
```

(\*) User\_Cycle for total run through Dhrystone with loops 5000000:  
1620000031  
So the DMIPS/MHz can be calculated by:  
 $1000000 / (\text{User\_Cycle} / \text{Number\_Of\_Runs}) / 1757 = 1.756642 \text{ DMIPS/MHz}$

**coremark**

```
CPU Frequency 108254227 Hz
Start to run coremark for 5000 iterations
2K performance run parameters for coremark.
CoreMark Size : 666
Total ticks : 1329373412
Total time (secs): 12.280106
Iterations/Sec : 407.162599
Iterations : 5000
Compiler version : GCC9.2.0
Compiler flags : -O2 -flto -funroll-all-loops -finline-limit=600 -ftree-dominator-opts -fno-if-conversion2 -fselective-scheduling -fno-code-hoisting -fno-common -funroll-loops -finline-functions -falign-functions=4 -falign-jumps=4 -falign-loops=4
Memory location : STACK
seedcrc : 0xe9f5
[0]crclist : 0xe714
[0]crcmatrix : 0x1fd7
[0]crcstate : 0x8e3a
[0]crcfinal : 0xbd59
Correct operation validated. See readme.txt for run and reporting rules.
CoreMark 1.0 : 407.162599 / GCC9.2.0 -O2 -flto -funroll-all-loops -finline-limit=600 -ftree-dominator-opts -fno-if-conversion2 -fselective-scheduling -fno-code-hoisting -fno-common -funroll-loops -finline-functions -falign-functions=4 -falign-jumps=4 -falign-loops=4 / STACK

Print Personal Added Additional Info to Easy Visual Analysis

(Iterations is: 5000
(total_ticks is: 1329373412
(*) Assume the core running at 1 MHz
So the CoreMark/MHz can be calculated by:
(Iterations*1000000/total_ticks) = 3.761170 CoreMark/MHz
```

**whetstone**

```
CPU Frequency 108254227 Hz
#####
Single Precision C Whetstone Benchmark Opt 3 32 Bit
Calibrate
  1.36 Seconds      1 Passes (x 100)
  6.78 Seconds      5 Passes (x 100)

Use 7 passes (x 100)

Single Precision C/C++ Whetstone Benchmark
```


Loop content	Result	MFLOPS	MOPS	Seconds
N1 floating point	-1.12475013732910156	1.518		0.089
N2 floating point	-1.12274742126464844	1.519		0.619
N3 if then else	1.00000000000000000		379846.656	0.000
N4 fixed point	12.00000000000000000		7.446	0.296
N5 sin,cos etc.	0.49909299612045288		0.157	3.706
N6 floating point	0.99999982118606567	1.487		2.539
N7 assignments	3.00000000000000000		562.837	0.002
N8 exp,sqrt etc.	0.75110614299774170		0.116	2.245
MWIPS		7.371		9.497
MWIPS/MHz		0.068		9.497

	dhystone	coremark	whetstone
QEMU benchmark	1.756642 DMIPS/MHz	3.761170 CoreMark/MHz	0.068MWIPS/MHz
SDK/app.rst	1.273270 DMIPS/MHz	3.081076 CoreMark/MHz	0.046MWIPS/MHz



# HummingBird Board

- a: FPGA\_RESET
- b: FPGA\_PROG
- c: MCU\_WKUP
- d: MCU\_RESET
- e1: User button 1
- e2: User button 2
- e3: User button header
- Y1: GCLK
- Y2: RTC\_CLK
- I: MCU\_FLASH
- 2: FPGA\_FLASH
- 3: FPGA\_JTAG
- 4: MCU\_JTAG
- 5: Power switch



蜂鸟FPGA评估板

## 蜂鸟FPGA评估板

主芯片FPGA型号为Xilinx XC7A100T

板载双晶振设计：100MHz主时钟和32.768KHz RTC时钟

配备两颗MT41K128M16JT-125K DDR III 颗粒

备独立的MCU\_FLASH芯片，此Flash用于存储RISC-V内核运行的程序文件

配备独立的FPGA\_Flash芯片，此Flash用于存储mcs格式的比特流文件

配备用户LEDs、按键及拨码开关

配备独立的32个MCU GPIO

配备多达126个引出的FPGA GPIO，用于用户自定义使用

配备了FPGA\_GPIO电平选择跳线接口，可以选择1.8V\2.5V\3.3V接口电平

# HummingBird Board

```
sifive_u_machine_init
```

```
create_fdt
```

```
riscv_find_and_load_firmware
```

```
riscv_load_initrd
```

```
cpus {  
    #address-cells = <1>;  
    #size-cells = <0>;  
    timebase-frequency = <32768>;  
    cpu0: cpu@0 {  
        device_type = "cpu";  
        reg = <0>;  
        status = "okay";  
        compatible = "riscv";  
        riscv,isa = "rv64imac";  
        mmu-type = "riscv,sv39";  
        clock-frequency = <8000000>;  
        cpu0_intc: interrupt-controller {  
            #interrupt-cells = <1>;  
            interrupt-controller;  
            compatible = "riscv,cpu-intc";  
        };  
    };  
};
```

```
qemu_fdt_add_subnode(fdt, "/cpus");  
qemu_fdt_setprop_cell(fdt, "/cpus", "timebase-frequency",  
    SIFIVE_CLINT_TIMEBASE_FREQ);  
qemu_fdt_setprop_cell(fdt, "/cpus", "#size-cells", 0x0);  
qemu_fdt_setprop_cell(fdt, "/cpus", "#address-cells", 0x1);  
  
for (cpu = ms->smp.cpus - 1; cpu >= 0; cpu--) {  
    int cpu_phandle = phandle++;  
    nodename = g_strdup_printf("/cpus/cpu%d", cpu);  
    char *intc = g_strdup_printf("/cpus/cpu%d/interrupt-  
controller", cpu);  
    char *isa;  
    qemu_fdt_add_subnode(fdt, nodename);  
    /* cpu 0 is the management hart that does not have mmu */  
    if (cpu != 0) {  
        qemu_fdt_setprop_string(fdt, nodename, "mmu-  
type", "riscv,sv39");  
        isa = riscv_isa_string(&s->soc.u_cpus.harts[cpu - 1]);  
    } else {  
        isa = riscv_isa_string(&s->soc.e_cpus.harts[0]);  
    }  
    qemu_fdt_setprop_string(fdt, nodename, "riscv,isa", isa);  
    qemu_fdt_setprop_string(fdt, nodename, "compatible", "riscv");  
    qemu_fdt_setprop_string(fdt, nodename, "status", "okay");  
    qemu_fdt_setprop_cell(fdt, nodename, "reg", cpu);  
    qemu_fdt_setprop_string(fdt, nodename, "device_type", "cpu");  
    qemu_fdt_add_subnode(fdt, intc);  
    qemu_fdt_setprop_cell(fdt, intc, "phandle", cpu_phandle);  
    qemu_fdt_setprop_string(fdt, intc, "compatible", "riscv,cpu-intc");  
    qemu_fdt_setprop(fdt, intc, "interrupt-controller", NULL, 0);  
    qemu_fdt_setprop_cell(fdt, intc, "#interrupt-cells", 1);  
    g_free(isa);  
    g_free(intc);  
    g_free(nodename);  
}
```

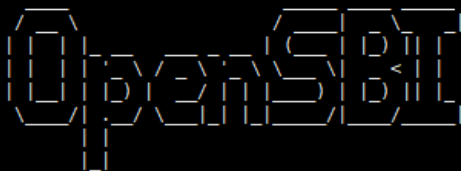


# HummingBird Board

```
$qemu-system-riscv64 \  
-M nuclei_u \  
-nographic \  
-m 1G \  
-bios none \  
-icount shift=0 \  
-append "root=/dev/ram rw console=ttyS0  
earlycon=sbi" \  
-kernel ~/workroom/risc-v/nuclei/nuclei-linux-  
sdk/work/opensbi/platform/nuclei/ux600/firm  
ware/fw_payload.elf \  
-initrd ./rootfs_riscv64.img
```

```
wang@lelouch:~/workroom/qemu$ ./run_hbird_linux.sh
```

```
OpenSBI v0.7
```

The logo for OpenSBI, featuring the text "OpenSBI" in a stylized, blocky font where the letters are interconnected.

```
Platform Name      : Nuclei UX600  
Platform Features  : timer,mfdeleg  
Platform HART Count : 1  
Boot HART ID       : 0  
Boot HART ISA      : rv64imafdcu  
BOOT HART Features : pmp,scounteren,mcounteren  
BOOT HART PMP Count : 16  
Firmware Base      : 0xa0000000  
Firmware Size      : 76 KB  
Runtime SBI Version : 0.2
```

```
MIDELEG : 0x0000000000000222  
MEDELEG : 0x000000000000b109  
PMP0     : 0x00000000a0000000-0x00000000a001ffff (A)  
PMP1     : 0x0000000000000000-0xffffffffffffff (A,R,W,X)  
[ 0.000000] OF: fdt: Ignoring memory range 0xa0000000 - 0xa0200000  
[ 0.000000] Linux version 5.7.0-13090-gad29b1fc8e7b (wang@lelouch) (riscv-nuclei-linux-gnu-gcc (GCC) 9.2.0, GNU ld (GNU Binutils) 2.32) #3 Mon Nov 23 14:11:23 CST 2020  
[ 0.000000] initrd not found or empty - disabling initrd  
[ 0.000000] Zone ranges:  
[ 0.000000] DMA32 [mem 0x00000000a0200000-0x00000000affffffff]  
[ 0.000000] Normal empty  
[ 0.000000] Movable zone start for each node  
[ 0.000000] Early memory node ranges  
[ 0.000000] node 0: [mem 0x00000000a0200000-0x00000000affffffff]  
[ 0.000000] Initmem setup node 0 [mem 0x00000000a0200000-0x00000000affffffff]  
[ 0.000000] software IO TLB: mapped [mem 0xabc7c000-0xafc7c000] (64MB)  
[ 0.000000] SBI specification v0.2 detected  
[ 0.000000] SBI implementation ID=0x1 Version=0x7  
[ 0.000000] SBI v0.2 TIME extension detected  
[ 0.000000] SBI v0.2 IPI extension detected  
[ 0.000000] SBI v0.2 RFENCE extension detected  
[ 0.000000] riscv: ISA extensions acim  
[ 0.000000] riscv: ELF capabilities acim
```

# Future Work

GitHub: <https://github.com/isrc-cas/plct-qemu>

- 代码版本更新和上游推送
- UX600系列, UX900系列支持
- SPI,ADC,QSPI,DMA等必要外设实现
- CLINT,PLIC,ECLIC等配置模块切换
- P-Extensions实现
- .....

欢迎实习生加入(。・∀・)ノ

<https://zhuanlan.zhihu.com/p/271625260>

# 谢谢

wangjunqiang@iscas.ac.cn